

Main Training

All models trained with:

batch_size = 4

epochs = 2

learning_rate = 0.001

momentum = 0.9

Mini-batch Gradient Descent

Cross Entropy Loss

Modifications made to models to fit 3x32x32 input size

Training accuracy given for all models

LeNet()

```
Model size:    parameters: 62006
Training time: 200.56246376037598
```

```
Accuracy of the network on the 10000 test images: 61 %
Accuracy for class: plane is 51.9 %
Accuracy for class: car   is 75.0 %
Accuracy for class: bird  is 55.6 %
Accuracy for class: cat   is 30.2 %
Accuracy for class: deer  is 47.7 %
Accuracy for class: dog   is 54.9 %
Accuracy for class: frog  is 67.3 %
Accuracy for class: horse is 74.3 %
Accuracy for class: ship  is 78.8 %
Accuracy for class: truck is 74.5 %
```

AlexNet()

```
Model size:    parameters: 36360074
Training time: 249.5857174396515
```

```
Accuracy of the network on the 10000 test images: 64 %
Accuracy for class: plane is 70.5 %
Accuracy for class: car is 74.0 %
Accuracy for class: bird is 38.6 %
Accuracy for class: cat is 59.2 %
Accuracy for class: deer is 45.5 %
Accuracy for class: dog is 56.6 %
Accuracy for class: frog is 74.7 %
Accuracy for class: horse is 67.9 %
Accuracy for class: ship is 78.5 %
Accuracy for class: truck is 78.2 %
```

Resnet-18()

```
Model size: parameters: 11184650
Training time: 259.30410146713257
```

```
Accuracy of the network on the 10000 test images: 52 %
Accuracy for class: plane is 40.7 %
Accuracy for class: car is 59.9 %
Accuracy for class: bird is 55.4 %
Accuracy for class: cat is 51.6 %
Accuracy for class: deer is 43.9 %
Accuracy for class: dog is 39.6 %
Accuracy for class: frog is 56.0 %
Accuracy for class: horse is 50.9 %
Accuracy for class: ship is 68.7 %
Accuracy for class: truck is 61.4 %
```

K-fold Cross Validation

K-fold tests were performed with:

batch_size = 40

folds = 5

other factors same as regular training

LeNet()

```
Fold 1
-----
Test set: Average loss: 0.0098, Accuracy: 2931/(10000.0) (29%)

Fold 2
-----
Test set: Average loss: 0.0096, Accuracy: 2964/(10000.0) (30%)

Fold 3
-----
Test set: Average loss: 0.0095, Accuracy: 3086/(10000.0) (31%)

Fold 4
-----
Test set: Average loss: 0.0094, Accuracy: 3157/(10000.0) (32%)

Fold 5
-----
Test set: Average loss: 0.0093, Accuracy: 3249/(10000.0) (32%)
```

AlexNet()

```
Fold 1
-----
Test set: Average loss: 0.0115, Accuracy: 1678/(10000.0) (17%)

Fold 2
-----
Test set: Average loss: 0.0114, Accuracy: 1613/(10000.0) (16%)

Fold 3
-----
Test set: Average loss: 0.0115, Accuracy: 2077/(10000.0) (21%)

Fold 4
-----
Test set: Average loss: 0.0115, Accuracy: 1667/(10000.0) (17%)

Fold 5
-----
Test set: Average loss: 0.0114, Accuracy: 1991/(10000.0) (20%)
```

ResNet-18()

Fold 1

Test set: Average loss: 0.0052, Accuracy: 6351/(10000.0) (64%)

Fold 2

Test set: Average loss: 0.0051, Accuracy: 6416/(10000.0) (64%)

Fold 3

Test set: Average loss: 0.0052, Accuracy: 6334/(10000.0) (63%)

Fold 4

Test set: Average loss: 0.0054, Accuracy: 6176/(10000.0) (62%)

Fold 5

Test set: Average loss: 0.0050, Accuracy: 6454/(10000.0) (65%)

ResNet()-18 was the best


Confusion matrix of ResNet-18:

```

[[407  11 300  34  52   5  12  14 129  36]
 [ 40 599  12  27  11   4  21   6  59 221]
 [ 22  11 554 120 103 106  41  14  17  12]
 [   4  15 128 516  80 120  60  19  26  32]
 [ 15   2 276 112 439  33  58  37  22   6]
 [   5   5 128 344  51 396  23  20  16  12]
 [   1   9  97 176 119  10 560   5  10  13]
 [   5   3  71  76 139 154  14 509   8  21]
 [ 81  30  90  44  18   2   7   1 687  40]
 [ 45 160  15  43  15   4  15  18  71 614]]

```

ResNet-18 Confusion Matrix



Three hypotheses

1. Increasing batch size of LeNet()

This will decrease training time, with lower test accuracy as a result. There will be fewer changes to model parameters, preventing it from accurately fitting the data. However, the smaller amount of gradient calculations will decrease training time.

```
Model size:    parameters: 62006
Training time: 39.29397535324097
```

```
Accuracy of the network on the 10000 test images: 10 %
Accuracy for class: plane is 0.0 %
Accuracy for class: car   is 0.0 %
Accuracy for class: bird  is 0.0 %
Accuracy for class: cat   is 0.0 %
Accuracy for class: deer  is 13.0 %
Accuracy for class: dog   is 0.0 %
Accuracy for class: frog  is 0.0 %
Accuracy for class: horse is 1.7 %
Accuracy for class: ship  is 89.0 %
Accuracy for class: truck is 0.0 %
```

An improvement in training time and decrease in accuracy is seen compared to the earlier result for LeNet(). The hypothesis is true.

2. Decreasing training input of LeNet() to 30.000

This will decrease accuracy due to overfitting of the smaller training set.

```
Accuracy of the network on the 10000 test images: 51 %
Accuracy for class: plane is 67.1 %
Accuracy for class: car   is 63.7 %
Accuracy for class: bird  is 22.8 %
Accuracy for class: cat   is 30.3 %
Accuracy for class: deer  is 51.9 %
Accuracy for class: dog   is 40.3 %
Accuracy for class: frog  is 53.0 %
Accuracy for class: horse is 61.5 %
Accuracy for class: ship  is 58.5 %
Accuracy for class: truck is 64.2 %
```

Accuracy did decrease with a smaller training input size (compare to earlier result of 61% with 50,000 input size).

3. Modifying LeNet()

Modifying LeNet() by adding two new fully connected layers as shown:

```
self.fc1 = nn.Linear(16*5*5, 120)
self.fcNew1 = nn.Linear(120, 100) # for hypothesis test (not in original model)
self.fc2 = nn.Linear(100, 84)
self.fcNew2 = nn.Linear(84, 32) # for hypothesis test
self.fc3 = nn.Linear(32, 10)
```

This will improve performance, since a larger number of parameters will allow the model to better fit the data.

```
Accuracy of the network on the 10000 test images: 49 %
Accuracy for class: plane is 60.5 %
Accuracy for class: car is 71.5 %
Accuracy for class: bird is 19.4 %
Accuracy for class: cat is 42.6 %
Accuracy for class: deer is 34.7 %
Accuracy for class: dog is 31.7 %
Accuracy for class: frog is 70.0 %
Accuracy for class: horse is 61.7 %
Accuracy for class: ship is 55.2 %
Accuracy for class: truck is 50.9 %
```

Performance was actually worse with increased model complexity. Too much information may be discarded in the smaller layers.

Notebook code (also available at:

<https://colab.research.google.com/drive/1b5e0wix6CxNkA0N08cvFSChJdfgo19jl?usp=sharing>

```
# -*- coding: utf-8 -*-
```

```
"""Helhoski_Assignment_3
```

Automatically generated by Colab.

Original file is located at

```
https://colab.research.google.com/drive/1b5e0wix6CxNkA0N08cvFSChJdfgo19jl
"""
```

```
#@title imports
```

```
import torch
```

```
import torchvision
```

```
import torchvision.transforms as transforms
```

```
import torch.optim as optim
```

```
from sklearn.model_selection import KFold
```

```
import time
```

```
import torch.cuda as cuda
```

```
#@title Loading Dataset
```

```
# much of model setup and training from:
```

```
https://pytorch.org/tutorials/beginner/blitz/cifar10\_tutorial.html
```

```
transform = transforms.Compose(
```

```
    [transforms.ToTensor(),
```

```
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

```
batch_size = 40
```

```
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
```

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
```

```
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                         shuffle=False, num_workers=2)
```

```
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

```
#@title training
```

```
def train(model, train_loader, optimizer):
```

```
    model.train()
```

```
    for batch_idx, (data, target) in enumerate(train_loader):
```



```
optimizer.zero_grad(set_to_none=True)
```

```
data = data.to('cuda')
target = target.to('cuda')
output = model(data)
output = output.to('cuda')
criterion = nn.CrossEntropyLoss()
loss = criterion(output, target)
loss.backward()
optimizer.step()
```

```
#@title Evaluation
```

```
def evalTotalSet(model):
```

```
    correct = 0
```

```
    total = 0
```

```
    model.eval()
```

```
    model.to('cuda')
```

```
    # since we're not training, we don't need to calculate the gradients for our outputs
    with torch.no_grad():
```

```
        for data in testloader:
```

```
            images, labels = data
```

```
            images = images.to('cuda')
```

```
            labels = labels.to('cuda')
```

```
            # calculate outputs by running images through the network
```

```
            outputs = model(images)
```

```
            # the class with the highest energy is what we choose as prediction
```

```
            _, predicted = torch.max(outputs.data, 1)
```

```
            total += labels.size(0)
```

```
            correct += (predicted == labels).sum().item()
```

```
print(f'Accuracy of the network on the 10000 test images: {100 * correct // total} %')
```

```
def evalCategories(model):
```

```
    # prepare to count predictions for each class
```

```
    correct_pred = {classname: 0 for classname in classes}
```

```
    total_pred = {classname: 0 for classname in classes}
```

```
    model.eval()
```

```
    model.to('cuda')
```

```
    # again no gradients needed
```

```
    with torch.no_grad():
```

```
        for data in testloader:
```

```
            images, labels = data
```

```
            images = images.to('cuda')
```

```
            labels = labels.to('cuda')
```

```
            outputs = model(images)
```

```
            _, predictions = torch.max(outputs, 1)
```

```
            # collect the correct predictions for each class
```

```

    for label, prediction in zip(labels, predictions):
        if label == prediction:
            correct_pred[classes[label]] += 1
            total_pred[classes[label]] += 1
# print accuracy for each class
for classname, correct_count in correct_pred.items():
    accuracy = 100 * float(correct_count) / total_pred[classname]
    print(f'Accuracy for class: {classname:5s} is {accuracy:.1f} %')

```

#@title Models

```

import torch.nn as nn
import torch.nn.functional as F

```

```

class LeNet(nn.Module): # variation of LeNet from
https://pytorch.org/tutorials/beginner/blitz/cifar10\_tutorial.html
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2,2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16*5*5, 120)
        self.fcNew1 = nn.Linear(120, 100) # for hypothesis test (not in original model)
        self.fc2 = nn.Linear(100, 84)
        self.fcNew2 = nn.Linear(84, 32) # for hypothesis test
        self.fc3 = nn.Linear(32, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x,1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fcNew1(x))
        x = F.relu(self.fc2(x))
        x = F.relu(self.fcNew2(x))
        x = self.fc3(x)
        return x

```

```

class AlexNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 96, 3)
        self.conv2 = nn.Conv2d(96, 256, 4, padding=1)
        self.conv3 = nn.Conv2d(256, 384, 3, padding=1)
        self.conv4 = nn.Conv2d(384, 384, 2, padding=1)
        self.conv5 = nn.Conv2d(384, 256, 3, padding=1)

```

```

self.pool = nn.MaxPool2d(2, 2)
self.fc1 = nn.Linear(256*4*4, 4096)
self.fc2 = nn.Linear(4096, 4096)
self.fc3 = nn.Linear(4096, 10)

```

```

def forward(self, x):
    x = self.pool(F.relu(self.conv1(x)))
    x = self.pool(F.relu(self.conv2(x)))
    x = F.relu(self.conv3(x))
    x = F.relu(self.conv4(x))
    x = self.pool(F.relu(self.conv5(x)))
    x = torch.flatten(x, 1) # flatten all dimensions except batch
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = self.fc3(x)
    return x

```

`#@title Residual`

```

class Residual(nn.Module):
    def __init__(self, input_channels, num_channels, use_1x1conv=False, strides=1):
        super().__init__()
        self.conv1 = nn.Conv2d(input_channels, num_channels, kernel_size=3, padding=1,
stride=strides)
        self.conv2 = nn.Conv2d(num_channels, num_channels, kernel_size=3, padding=1)

        if use_1x1conv:
            self.conv3 = nn.Conv2d(input_channels, num_channels, kernel_size=1, stride=strides)
        else:
            self.conv3 = None

        self.bn1 = nn.BatchNorm2d(num_channels)
        self.bn2 = nn.BatchNorm2d(num_channels)

```

```

def forward(self, X):
    Y = F.relu(self.bn1(self.conv1(X)))
    Y = self.bn2(self.conv2(Y))
    if self.conv3:
        X = self.conv3(X)
    Y += X
    return F.relu(Y)

```

`#@title ResNet-18`

```

b1 = nn.Sequential(nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3),
                    nn.BatchNorm2d(64), nn.ReLU(),
                    nn.MaxPool2d(kernel_size=3, stride=2, padding=1))

```

```

def resnet_block(input_channels, num_channels, num_residuals, first_block=False):

```

```

blk = []
for i in range(num_residuals):
    if i == 0 and not first_block:
        blk.append(Residual(input_channels, num_channels, use_1x1conv=True, strides=2))
    else:
        blk.append(Residual(num_channels, num_channels))
return blk

b2 = nn.Sequential(*resnet_block(64, 64, 2, first_block=True))
b3 = nn.Sequential(*resnet_block(64, 128, 2))
b4 = nn.Sequential(*resnet_block(128, 256, 2))
b5 = nn.Sequential(*resnet_block(256, 512, 2))

resnet18Inst = nn.Sequential(b1, b2, b3, b4, b5,
                             nn.AdaptiveAvgPool2d((1,1)),
                             nn.Flatten(), nn.Linear(512, 10))

#@title training resnet

# training all 3 models and evaluating on test set
resnet18Inst.to('cuda')

batch_size = 4
epochs = 2
optimizer = optim.SGD(resnet18Inst.parameters(), lr=0.001, momentum=0.9)
# criterion = F.nll_loss()

trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True)

# method of counting parameters from:
https://discuss.pytorch.org/t/how-do-i-check-the-number-of-parameters-of-a-model/4325
print("Model size:  parameters:", sum(p.numel() for p in resnet18Inst.parameters() if
p.requires_grad))
start = time.time()
for epoch in range(epochs):
    train(resnet18Inst, trainloader, optimizer)
end = time.time()
print("Training time: ", end-start)

import numpy
import matplotlib.pyplot as plt

resnet18Inst = nn.Sequential(b1, b2, b3, b4, b5,
                             nn.AdaptiveAvgPool2d((1,1)),
                             nn.Flatten(), nn.Linear(512, 10))
PATH = './cifar_resNet.pth'

```

```
resnet18Inst.load_state_dict(torch.load(PATH, weights_only=True))
```

```
c_matrix = numpy.zeros((10,10), dtype=int)
evalCategories(resnet18Inst)
print(c_matrix)
```

#some code from:

<https://stackoverflow.com/questions/20998083/show-the-values-in-the-grid-using-matplotlib>

```
plt.matshow(c_matrix)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('ResNet-18 Confusion Matrix')
plt.show()
```

```
#@title training lenet
# training all 3 models and evaluating on test set
leNetInst = LeNet()
leNetInst.to('cuda')
```

```
batch_size = 4
epochs = 2
optimizer = optim.SGD(leNetInst.parameters(), lr=0.001, momentum=0.9)
#criterion = F.nll_loss()
```

```
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True)
```

```
# method of counting parameters from:
https://discuss.pytorch.org/t/how-do-i-check-the-number-of-parameters-of-a-model/4325
print("Model size:  parameters:", sum(p.numel() for p in leNetInst.parameters() if
p.requires_grad))
start = time.time()
for epoch in range(epochs):
    train(leNetInst, trainloader, optimizer)
end = time.time()
print("Training time: ", end-start)
```

```
evalTotalSet(leNetInst)
evalCategories(leNetInst)
```

```
# training all 3 models and evaluating on test set
alexInst = AlexNet()
alexInst.to('cuda')
```

```

batch_size = 4
epochs = 2
optimizer = optim.SGD(alexInst.parameters(), lr=0.001, momentum=0.9)
# criterion = F.nll_loss()

trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=2)
# method of counting parameters from:
https://discuss.pytorch.org/t/how-do-i-check-the-number-of-parameters-of-a-model/4325
print("Model size:  parameters:", sum(p.numel() for p in alexInst.parameters() if
p.requires_grad))
start = time.time()
for epoch in range(epochs):
    train(alexInst, trainloader, optimizer)
end = time.time()
print("Training time: ", end-start)

PATH = './cifar_alexNet.pth'
torch.save(alexInst.state_dict(), PATH)

#@title k-fold
def run_k_fold(folds, batch_size, epochs):

    # kfold method/some code from:
https://saturncloud.io/blog/how-to-use-kfold-cross-validation-with-dataloaders-in-pytorch/#step-1
    -importing-the-required-libraries
    kf = KFold(n_splits=folds)

    # Loop through each fold
    for fold, (train_idx, test_idx) in enumerate(kf.split(trainset)):
        print(f"Fold {fold + 1}")
        print("-----")

        # Define the data loaders for the current fold
        train_loader = torch.utils.data.DataLoader(dataset=trainset, batch_size=batch_size,
                                                    sampler=torch.utils.data.SubsetRandomSampler(train_idx),
                                                    pin_memory=True, num_workers=2)
        test_loader = torch.utils.data.DataLoader(dataset=trainset, batch_size=batch_size,
                                                    sampler=torch.utils.data.SubsetRandomSampler(test_idx),
                                                    pin_memory=True, num_workers=2)

        # have to redefined entire model inside block so that residual blocks are reset
        b1 = nn.Sequential(nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3),
                            nn.BatchNorm2d(64), nn.ReLU(),
                            nn.MaxPool2d(kernel_size=3, stride=2, padding=1))

        def resnet_block(input_channels, num_channels, num_residuals, first_block=False):

```

```

blk = []
for i in range(num_residuals):
    if i == 0 and not first_block:
        blk.append(Residual(input_channels, num_channels, use_1x1conv=True, strides=2))
    else:
        blk.append(Residual(num_channels, num_channels))
return blk

b2 = nn.Sequential(*resnet_block(64, 64, 2, first_block=True))
b3 = nn.Sequential(*resnet_block(64, 128, 2))
b4 = nn.Sequential(*resnet_block(128, 256, 2))
b5 = nn.Sequential(*resnet_block(256, 512, 2))

model = nn.Sequential(b1, b2, b3, b4, b5,
                      nn.AdaptiveAvgPool2d((1,1)),
                      nn.Flatten(), nn.Linear(512, 10))
model.to('cuda')
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

for epoch in range(epochs):
    train(model, train_loader, optimizer)

model.eval()
test_loss = 0
correct = 0
with torch.no_grad():
    for data, target in test_loader:
        data = data.to('cuda')
        target = target.to('cuda')
        output = model(data)
        criterion = nn.CrossEntropyLoss()
        test_loss += criterion(output, target).item()
        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

test_loss /= len(test_loader.dataset)
accuracy = 100.0 * correct / (len(test_loader.dataset)/folds)
print(f"Test set: Average loss: {test_loss:.4f}, Accuracy:
{correct}/{(len(test_loader.dataset)/folds)} ({accuracy:.0f}%)")
print()

```